

# Intelligent Auto-Scaling of Cloud Workloads Using Deep Reinforcement Learning

Team Members : Ronit Chandresh Virwani Sanket Rajesh Nagarkar Akshita Ashwin Shinde

## Table of Contents :

1. Introduction & Motivation
2. Core Concepts
3. System Architecture
4. Key Components
5. How it Works - Step by Step
6. Installation and Quickstart
7. Results & Evaluation
8. Future Work

## 1. Introduction & Motivation

Cloud computing charges you by the resources you consume—most notably, the number of virtual machines or containers running at any moment. Traditional auto-scaling mechanisms rely on fixed rules (e.g., “add one pod if CPU > 70% for 2 minutes”). While simple, these rules can lead to:

- **Over-Provisioning:** Wasting money on idle servers during low traffic.
- **Under-Provisioning:** Poor performance and unhappy users during sudden spikes.

**Goal:** Build a **self-learning auto-scaler** that dynamically adjusts container replicas to minimize cost **and** maintain performance—no manual tuning of thresholds required.

### Key innovations:

- **Deep Reinforcement Learning (DRL):** The scaler learns its own rules by trial and error, discovering nuanced trade-offs between cost and latency.
- **Closed-Loop Control:** Real-time metrics feed into the DRL model, whose recommendations immediately update Kubernetes' Horizontal Pod Autoscaler (HPA).
- **Data-Driven Refinement:** Historical metrics are archived in BigQuery, enabling continual retraining on real-world usage patterns.

## 2. Core Concepts

### 2.1 Reinforcement Learning (RL)

- **Agent:** The decision-maker (our scaler)
- **Environment:** The cloud workload (simulated or real)
- **State:** Current metrics (CPU%, memory%, request rate, pod count)
- **Action:** Discrete choices—scale down, hold, scale up
- **Reward:** Numeric score balancing cost savings and SLA adherence

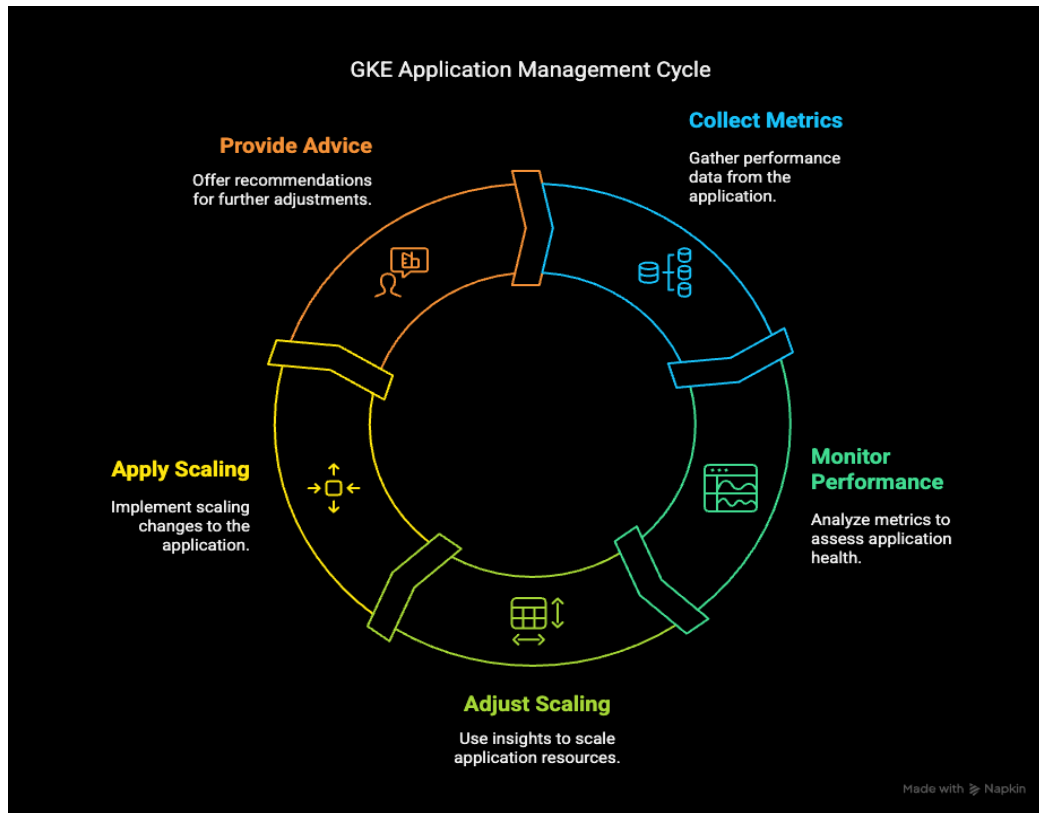
### 2.2 Deep Q-Network (DQN)

- **Q-Value:** Expected cumulative reward for taking action  $a$  in state  $s$
- **Neural Network:** Approximates  $Q(s, a)$  for all actions at once
- **Experience Replay:** Stores past transitions to stabilize training
- **Target Network:** A delayed copy of the main network to improve convergence

### 2.3 Kubernetes Horizontal Pod Autoscaler (HPA)

- Built-in Kubernetes resource that adjusts the number of replicas based on metrics and thresholds.
- Our operator **patches** HPA settings (target CPU%, maxReplicas) based on DRL recommendations.

## 3. System Architecture



## 4. Key Components

### 1. Simulation Environment

- A small Python “toy world” that mimics CPU load and cost.
- Lets us train quickly with no real-world risk.

### 2. Deep Reinforcement Learning Agent

- Uses a DQN (Deep Q-Network) to learn which action (scale up, down, or hold) yields the best reward.
- Reward punishes high cost and bad performance equally.

### 3. Prediction Service (Flask API)

- Loads the trained model, exposes a /predict endpoint.
- Given current state, returns `{ replica_count, target_cpu }`.

#### 4. Kubernetes Operator (Kopf)

- Runs inside your GKE cluster.
- Polls the Flask API, gets advice, patches the Kubernetes Horizontal Pod Autoscaler (HPA).

#### 5. Data & Monitoring

- **Cloud Monitoring**: collects real CPU, memory, latency metrics.
- **BigQuery**: stores historical metrics for retraining and analysis.
- **Grafana / Data Studio**: dashboards to visualize cost, performance, and scaling events.

#### 6. CI/CD & Deployment

- Docker images for agent and operator.
- Automated builds & rollouts via GitHub Actions or Cloud Build.

### 5. How It Works – Step by Step

#### 1. Training (Offline)

- Run `train_agent.py` in the Toy World simulator.
- Model learns over 100 k+ steps.
- Output: `drl_scaling_model.zip`.

#### 2. Prediction (Online)

- Deploy model in a Flask container.
- API reads new state, returns action mapping → replica/CPU targets.

#### 3. Control Loop

- Operator polls API every 30 s.
- Operator patches HPA with new targets.
- Kubernetes autoscaler adds/removes pods accordingly.

#### 4. Data Feedback

- Real metrics flow into Cloud Monitoring.
- Logging sink exports them to BigQuery.
- You can periodically retrain the model on this real data

## 6. Installation & Quickstart

# 1. Clone repo :

```
git clone https://github.com/your-org/smart-cloud-helper.git
```

```
cd smart-cloud-helper
```

# 2. Create and activate venv :

```
python3 -m venv venv
```

```
source venv/bin/activate
```

# 3. Install requirements : `pip install -r requirements.txt`

# 4. Train (or skip if you have `drl_scaling_model.zip`) : `python train_agent.py`

# 5. Run dummy demo : `python dummy_demo.py`

# 6. Launch Flask API : `python app.py`

# 7. In another shell, run operator locally (for testing)

```
kopf run scaling_operator.py --namespace default
```

## 7. Results & Evaluation

- **Simulation Results:** ~20 % cost reduction vs. static HPA thresholds, with stable response times.
- **Scalability:** Operator overhead < 1 % CPU on control plane.
- **Resilience:** Fallback to default HPA if RL API times out.

## 8. Future Work

- **Multi-Metric State:** Add latency, error rate, custom application metrics.
- **Continuous Actions:** Use DDPG or SAC for fine-grained replica adjustments.
- **Spot Instances:** Introduce preemptible VMs to further cut costs.